

Simple linear regression - Exercise

You are given a real estate dataset.

Real estate is one of those examples that every regression course goes through as it is extremely easy to understand and there is a (almost always) certain causal relationship to be found.

The data is located in the file: 'real_estate_price_size.csv'.

You are expected to create a simple linear regression (similar to the one in the lecture), using the new data.

Apart from that, please:

- Create a scatter plot (with or without a regression line)
- Calculate the R-squared
- Display the intercept and coefficient(s)
- Using the model make a prediction about an apartment with size 750 sq.ft.

Note: In this exercise, the dependent variable is 'price', while the independent variable is 'size'.

Good luck!

Import the relevant libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

from sklearn.linear_model import LinearRegression
```

Load the data

```
In [2]: data = pd.read_csv('real_estate_price_size.csv')
```

```
In [3]: data.head()
```

```
Out[3]:
```

	price	size
0	234314.144	643.09
1	228581.528	656.22
2	281626.336	487.29
3	401255.608	1504.75
4	458674.256	1275.46

Create the regression

Declare the dependent and the independent variables

```
In [4]: y = data['price']  
x = data['size']
```

Explore the data

```
In [5]: x.shape
```

```
Out[5]: (100,)
```

```
In [6]: y.shape
```

```
Out[6]: (100,)
```

Transform the inputs into a matrix (2D object)

```
In [7]: x_matrix = x.values.reshape(-1,1)  
x_matrix.shape
```

```
Out[7]: (100, 1)
```

Regression itself

```
In [8]: reg = LinearRegression()  
reg.fit(x_matrix,y)
```

```
Out[8]: LinearRegression()
```

Calculate the R-squared

```
In [9]: reg.score(x_matrix,y)
```

```
Out[9]: 0.7447391865847586
```

Find the intercept

```
In [10]: reg.intercept_
```

```
Out[10]: 101912.60180122912
```

Find the coefficients

```
In [11]: reg.coef_
```

```
Out[11]: array([223.17874259])
```

Making predictions

You find an apartment online with a size of 750 sq.ft.

All else equal what should be its price according to the model?

```
In [12]: new_data = pd.DataFrame(data=[656.22, 1504.75], columns=['size'])
new_data
```

```
Out[12]:
```

	size
0	656.22
1	1504.75

```
In [13]: reg.predict(new_data)
```

```
Out[13]: array([248366.95626666, 437740.81472046])
```

```
In [14]: new_data['predicted_price'] = reg.predict(new_data)
new_data
```

```
Out[14]:
```

	size	predicted_price
0	656.22	248366.956267
1	1504.75	437740.814720

```
In [17]: plt.scatter(x,y)
yhat = reg.coef_ * x_matrix + reg.intercept_
fig = plt.plot(x,yhat, lw=4, c='orange', label='regression line')
plt.xlabel('size', fontsize=20)
plt.ylabel('price', fontsize=20)
plt.show()
```

